


GPT-5.4 F1 across decreasing session context

Lines connect four evaluated categorical horizons; they do not interpolate untested durations. *Author-supplied F1-only result.

Minimal (5) Mercury-style (20) Combined (25)

Balanced evaluation (50.00% malicious)



Deployment GPT subset (49.36% malicious)

